

1 Review of useful concepts and Introduction

1.1 Multivariate Gaussian

$$f(x) = \frac{1}{2\pi\sqrt{|\Sigma|}} e^{-\frac{1}{2}(x-\mu)^T \Sigma^{-1}(x-\mu)}$$

Suppose we have a Gaussian random vector $X_V \sim N(\mu_V, \Sigma_{VV})$.

Suppose we take two disjoint subsets of V : $A = i_1, \dots, i_k$ and $B = j_1, \dots, j_m$.

Then, the conditional distribution:

$P(X_A | X_B = x_B) = N(\mu_{A|B}, \Sigma_{A|B})$ is Gaussian:

$$\mu_{A|B} = \mu_A + \Sigma_{AB} \Sigma_{BB}^{-1} (x_B - \mu_B)$$

$$\Sigma_{A|B} = \Sigma_{AA} - \Sigma_{AB} \Sigma_{BB}^{-1} \Sigma_{BA}$$

1.2 Convex / Jensen's inequality

$g(x)$ is convex $\Leftrightarrow x_1, x_2 \in \mathbb{R}, \lambda \in [0, 1] : g''(x) > 0$
 $g(\lambda x_1 + (1 - \lambda)x_2) \leq \lambda g(x_1) + (1 - \lambda)g(x_2)$
 $\varphi(E[X]) \leq E[\varphi(X)]$

1.3 Review Probability

Probability space (Ω, F, P) : Set of atomic events Ω . Set of all non-atomic events $(\sigma\text{-Algebra})$. $F \in 2^\Omega$. Probability measure: $P : F \rightarrow [0, 1]$

Bayes' rule: $P(B|A) = P(A, B)/P(A) = P(A|B)P(B)/P(A)$, where $P(A) = \sum_b P(A|B)P(B)$

Union: $P(A \cup B) = P(A) + P(B) - P(A \cap B)$

Rules for joint distributions:

Sum rule (Marginalization):

$$P(X_{1:i-1}, X_{i+1:n}) = \sum_{x_i} P(X_{1:i-1}, X_i = x_i, X_{i+1:n})$$

Product rule (Chain rule):

$$P(X_{1:n}) = P(x_1)P(x_2|x_1) \dots P(x_n|x_{1:n-1})$$

Conditional Independence:

$$X \perp Y | Z \text{ iff } P(X, Y | Z) = P(X | Z)P(Y | Z)$$

$$\text{If } P(Y | Z) > 0 \Rightarrow P(X | Z, Y) = P(X | Z)$$

Properties of Conditional Independence:

Symmetry: $X \perp Y | Z \Rightarrow Y \perp X | Z$

$$\text{Decomposition: } X \perp (Y, W) | Z \Rightarrow X \perp Y | Z$$

$$\text{Contraction: } (X \perp Y | Z) \wedge (X \perp W | Y, Z) \Rightarrow X \perp Y, W | Z$$

$$\text{Weak union: } X \perp Y, W | Z \Rightarrow X \perp Y | Z, W$$

$$\text{Intersection: } (X \perp Y | W, Z) \wedge (X \perp W | Y, Z) \Rightarrow X \perp Y, W | Z$$

2 Bayesian Networks

2.1 Basic concepts

A Bayesian network (G, P) consists of:

- A BN structure G (directed, acyclic graph)
- A set of conditional probability distributions
- (G, P) defines the joint distribution:

$$P(X_1, \dots, X_n) = \prod_i P(X_i | pa_{X_i})$$

BNs with 3 nodes:

2.2 Active trails and d-separation

An undirected path in a BN structure G is called active trail for observed variables $O \in X_1, \dots, X_n$ if for every consecutive triple of variables X, Y, Z on the path:

- **indirect causal effect:**

$X \rightarrow Y \rightarrow Z$ and Y unobserved

- **direct evidential effect:**

$X \leftarrow Y \leftarrow Z$ and Y unobserved

- **common cause:**

$X \leftarrow Y \rightarrow Z$ and Y unobserved.

- **common effect:**

$X \rightarrow Y \leftarrow Z$ and Y or any of Y 's descendants is observed.

Any variables X_i and X_j for which there is no active trail for observations O are called d-separated by O .

Theorem: $d\text{-sep}(X_i; X_j | O) \Rightarrow X \perp Y | Z$

Converse does not hold in general!

3 Exact inference (tree-structured BN)

3.1 Variable elimination

- Given a BN and query $P(X | E = e)$

- Choose an ordering of X_1, \dots, X_n **Eliminate variables from the outside in!**

- Set up initial factors: $f_i = P(X_i | pa_{X_i})$

- For $i = 1 : n, X_i \notin X, E$

- Collect and multiply all factors f that include X_i

- Generate new factor by marginalizing out X_i : $g_{X_i} = \sum_{x_i} \prod_j f_j$

- Add g to set of factors

- Renormalize $P(x, e)$ to get $P(x | e)$

Variable elimination for polytrees:


- Pick a root, (avoiding X and E)

- Orient edges towards root

- Eliminate variables according to topological order

3.2 Avoiding recomputation: factor graphs

FG for a BN is a bipartite graph consisting of variables (circles) and factors (rectangles). **It is not a unique representation.**



3.2.1 Sum-product/Belief Propagation (BP) Algorithm:

- Initialize all messages as uniform distribution

- Until converged to:

- Pick a root in the factor graph and reorient the edges towards this root.

- Update messages according to this ordering. Do passes from leaves to root and from root to leaves.

- If a leaf node is a variable node: $\mu_{x \rightarrow f}(x) = 1$

- If a leaf node is a factor node: $\mu_{f \rightarrow x}(x) = f(x)$

- Messages from node v to factor u :

$$\mu_{v \rightarrow u}(x_v) = \prod_{u' \in N(v) \setminus \{u\}} \mu_{u' \rightarrow v}(x_{v'})$$

- Messages from factor u to node v :

$$\mu_{u \rightarrow v}(x_v) = \sum_{x_{u'} \sim x_v} f_u(x_u) \prod_{v' \in N(u) \setminus \{v\}} \mu_{v' \rightarrow u}(x_{v'})$$

- Break once all messages change by $\leq \epsilon$

Hope: after convergence, we have:

$$P(X_v = x_v) = \frac{1}{Z} \prod_{u \in N(v)} \mu_{u \rightarrow v}(x_v)$$

$$P(\vec{X}_u = \vec{x}_u) = \frac{1}{Z} f_u(\vec{x}_u) \prod_{v \in N(u)} \mu_{v \rightarrow u}(x_v)$$

If we have a polytree Bayesian network:

- Choose one node as root

- Send messages from leaves to root and from root to leaves

4 Approximate inference (loopy networks)

With loopy graphs, BP is often **overconfident/oscillates**.

4.1 Variable elimination for MPE (most probable explanation):

- Given BN and evidence $E=e$

- Choose an ordering of X_1, \dots, X_n

- Set up initial factors $f_i = P(X_i | pa_i)$

- For $i = 1 : n, X_i \notin E$:

- Collect and multiply all factors f_j that include X_i

- Generate new factor by maximizing out X_i :

$$g_i = \max_{x_i} \prod_j f_j$$

- Add g to set of factors

- For $i = n-1 : 1, X_i \notin E$: $\hat{x}_i = \text{argmax}_{x_i} g_i(x_i, \hat{x}_{i+1:n})$

Retrieving MAP from Max-Product (MAP = MPE for a subset of RVs):

- Define max-marginals:

$$P_{\max}(X_v = x_v) := \max_{x \sim x_v} P(x)$$

- For tree factor graphs, max-product computes max-marginals:

$$P_{\max}(X_v = x_v) \propto \prod_{u \in N(v)} \mu_{u \rightarrow v}(x_v)$$

- Can retrieve MAP solution from these (must be careful when ties need to be broken).

4.2 Sampling based inference: compute marginals as expectations

Hoeffding's inequality: Suppose f is bounded in $[0, C]$. Then:

$$P(|E_P[f(X)] - \frac{1}{N} \sum_{i=1}^N f(x_i)| \geq \epsilon) \leq 2 \exp\left(\frac{-2N\epsilon^2}{C^2}\right)$$

Monte Carlo Sampling from a BN:

- Sort variables in topological ordering X_1, X_n

- For $i = 1$ to n , sample:

$$x_i \sim P(X_i = x_1, \dots, X_{i-1} = x_{i-1})$$

Rejection Sampling:

- Collect samples over all variables:

$$\hat{P}(X_A = x + A | X_B = x_B) = \frac{\text{Count}(x_A, x_B)}{\text{Count}(x_B)}$$

- Throw away samples that disagree with x_B

- Count fraction of x_a on remaining samples

4.2.1 Directly sampling from the posterior: MCMC

Markov Chain:: A (stationary) MC is a sequence of RVs X_1, \dots, X_N , with prior $P(X_1)$ and transition probabilities $P(X_{t+1} | X_t)$ inde-

pendent of t .

Markov assumpt.: $X_{1:t-1} \perp X_{t+1:T} | X_t, \forall t > 1$

Stationarity assumption:

$$P(X_{t+1} = x | X_t = x') = P(X_t = x | X_{t-1} = x'), \forall t > 1$$

If ergodic (= there exists a finite t such that every state can be reached in exactly t steps), then: it has a unique and positive stationary distribution $\pi(X) > 0$, such for all x :

$$\lim_{N \rightarrow \infty} P(X_N = x) = \pi(x) \text{ and } \pi(X) \perp P(X_1).$$

If MC satisfies the **detailed balance equation** (for unnormalized distribution Q , for all x, x' : $Q(x)P(x'|x) = Q(x')P(x|x')$), then the MC has stationarity distribution $\pi(X) = 1/Z Q(X)$.

Designing Markov Chains:

- Proposal distribution $R(X'|X)$: given $X_t = x$, sample "proposal" $x' \sim R(X'|X = x)$

- Acceptance distribution

- Suppose $X_t = x$

- With probability $\alpha = \min\left\{1, \frac{Q(x')R(x|x')}{Q(x)R(x'|x)}\right\}$

set: $X_{t+1} = x'$

- With probability $1 - \alpha$, set $X_{t+1} = x$

MCMC for graphical models: Gibbs sampling (Random Vs Practical variant):

- Start with initial assignment x to all variables

- Fix observed variables $X_B = x_B$

- For $t = 1$ to ∞ , do:

- Pick a variable i uniformly at random from

$\{1, \dots, n\} \setminus B$ / **Set ordering**, and then, for each X_i (except those in B)

- Set $v_i =$ values of all x except x_i

- Sample x_i from $P(X_i | v_i)$

5 Dynamical models (include time)

5.1 Examples with one variable per time step

X_1, \dots, X_T (unobserved) hidden states

Y_1, \dots, Y_T (noisy) observations

HMMs (polytrees: can use belief propagation): X_i categorical, Y_i categorical (or arbitrary)

Kalman filters: X_i, Y_i Gaussian distributions

- $P(X_1)$: prior belief about location at time i

- $P(X_{t+1} | X_t)$: 'Motion model' (how do I expect my target to move in the environment?):

$$X_{t+1} = FX + \epsilon_t \text{ where } \epsilon_t \sim N(0, \Sigma_\epsilon)$$

- $P(Y_t | X_t)$: 'Sensor model' (what do I observe if target is at location X_t ?) $Y_t = HX_t + \eta_t$ where

$$\eta_t \sim N(0, \Sigma_\eta)$$

5.2 Inference tasks

Filtering: $P(X_t | y_{1:t})$ Is it raining today?

Prediction: $P(X_{t+\tau} | Y_{1:t})$ Rain 5 days from now?

Example for one step: $P(X_{t+1} | Y_{1:t}) = \sum_x P(X_{t+1}, X_t = x_t | Y_{1:t}) = \sum_x P(X_{t+1} | X_t = x_t) P(X_t | Y_{1:t})$ (with KFs, you need **integrals**!)

Smoothing: $P(X_t | y_{1:t})$ with $\tau < t$ Did it rain

last week.] Can use sum-product (aka forward-backward).]

MPE: $\text{argmax}_{x_{1:T}} P(x_{1:T}|y_{1:T})$ Can use max product (aka Viterbi algorithm).

Bayesian filtering: Start with $P(X_1)$:

At time t , assume we have $P(X_t|y_{1:t-1})$

Conditioning: $P(X_t|y_{1:t}) = \frac{P(X_t|y_{1:t-1})P(y_t|X_t)}{\sum_{x_t} P(X_t|y_{1:t-1})P(y_t|X_t)}$

Prediction ($O(n^2)$ vs $O(n)$ in conditioning):

$P(X_{t+1}|y_{1:t}) = \sum_{x_t} P(X_{t+1}|X_t)P(X_t|y_{1:t})$

Since HMM is a polytree, smoothing/MPE can be computed by VE/BP. Kalman filtering: Bayesian filtering for continuous problems. RV corrupted by Gaussian distributions with zero mean. Bayesian filtering is basically the same, except that sums turn to integrals. General Kalman update

- Transition model: $P(x_{t+1}|x_t) = N(x_{t+1}; Fx_t, \Sigma_x)$

- Sensor model: $P(y_t|x_t) = N(y_t; Hx_t, \Sigma_y)$

- Kalman update:

$\mu_{t+1} = F\mu_t + K_{t+1}(y_{t+1} - HF\mu_t)$

$\Sigma_{t+1} = (I - K_{t+1})(F\Sigma_t F^T + \Sigma_x)$

- Kalman gain: $K_{t+1} = (F\Sigma_t F^T + \Sigma_x)H^T(H(F\Sigma_t F^T + \Sigma_x)H^T + \Sigma_y)^{-1}$

5.3 Examples with > 1 variable per time step

Dynamic Bayesian Networks: a BN at every time step

These models typically have many loops. Exact inference is usually intractable.

5.4 Approx. infer. for filtering (DBNs and non-linear Kalman filters): Particle filtering

Suppose: $P(X_t|y_{1:t}) \approx \frac{1}{N} \sum_{i=1}^N \delta_{x_{i,t}}$, where δ is the indicator function. **Prediction:** Propagate each particle: $x'_i \sim P(X_{t+1}|x_{i,t})$

Conditioning:

- weight particles $w_i = \frac{1}{Z} P(y_{t+1}|x'_i)$

- resample N particles $x_{i,t+1} \sim \frac{1}{Z} \sum_{i=1}^N w_i \delta_{x'_i}$

Conclusion we came to: $Z = \sum_{i=1}^N w_i \delta_{x_i}$

6 Probabilistic Planning

6.1 Markov Decision Processes

An MDP is specified by a quintuple: $(X, A, r, P(x'|x, a), \gamma)$, where X are states, A are actions, $r(x, a)$ is a reward function and transition probabilities:

$P(x'|x, a) = \text{Prob}(\text{Next state} = x' | \text{Action } a)$

Objective: find a stationary policy $\pi : S \rightarrow A$ that maximizes the sum of cumulative rewards.

Value of a state given a policy: sum of cumulative rewards, given that the initial state is this state \rightarrow **Bellman equation:**

$$V^\pi(s) = \mathbb{E} \left[\sum_{t=0}^{\infty} \gamma^t r(s_t, \pi(s_t), s_{t+1}) | s_0 = s \right]$$

$$= \sum_{s' \in S} P(s'|s, \pi(s)) [r(s, \pi(s), s') + \gamma V^\pi(s')] \\ = r(s, \pi(s)) + \gamma \sum_{s' \in S} P(s'|s, \pi(s)) V^\pi(s')$$

Theorem (Bellman): a policy is optimal iff it is greedy w.r.t. its induced value function!

$V^*(x) = \max_a [r(x, a) + \gamma \sum_{x'} P(x'|x, a) V^*(x')]$

Bellman equation mais geral:

$V^*(x) = \max_a [\sum_{x'} P(x'|x, a) (r(a, x, x') + \gamma V^*(x'))]$

Optimal policy:

$\pi^*(s) = \text{argmax}_{a \in A} [r(x, a) + \gamma \sum_{x'} P(x'|x, a) V^*(x')]$

6.2 Policy iteration (Cost $O(S^3 + SAA)$)

Start with an arbitrary (e.g. random) policy π . Until converged, do:

- Compute value function $V^\pi(x)$

- Compute greedy policy π_G w.r.t. V^π

- Set $\pi \leftarrow \pi_G$

Guaranteed to monotonically improve and to converge to an optimal policy π^* in $O(n^2 m / (1 - \gamma))$ iterations (converges in polynomial number of iterations)!

6.3 Value iteration (Cost $O(SAA)$)

Initialize $V_0(x) = \max_a r(x, a)$

For $t = 1$ to ∞ :

- For each (x, a) , let:

$Q_t(x, a) = r(x, a) + \gamma \sum_{x'} P(x'|x, a) V_{t-1}(x')$

- For each x , let $V_t(x) = \max_a Q_t(x, a)$

- Break if $\|V_t - V_{t-1}\|_\infty = \max_x |V_t(x) - V_{t-1}(x)| \leq \epsilon$

Then choose greedy policy w.r.t. V_t .

Guaranteed to converge to ϵ -optimal policy (finds approximate solution in polynomial number of iterations)!

6.4 POMDP = Belief-state MDP

States = beliefs over states for original POMDP

$B = \Delta(1, \dots, n) = \{b : 1, \dots, n \rightarrow [0, 1], \sum_x b(x) = 1\}$

Actions: same as original MDP

Transition model:

- Stochastic observation:

$P(Y_t|b_t) = \sum_{x=1}^n P(Y_t|X_t = x) b_t(x)$

- State update (Bayesian filtering!), given b_t, y_t, a_t : $b_{t+1}(x') = \frac{1}{Z} \sum_x b_t(x) P(y_t|x) P(X_{t+1} = x'|X_t = x, a_t)$

Reward function: $r(b_t, a_t) = \sum_x b_t(x) r(x, a_t)$

6.5 Example of approx. solution to POMDPs: Policy gradients

- Assume parameterized policy: $\pi(b) = \pi(b; \theta)$

- For each parameter θ the policy induces a Markov chain

- Can compute expected reward $J(\theta)$ by sampling.

- Find optimal parameters through search (gradient ascent): $\theta^* = \text{argmax}_\theta J(\theta)$

7 Learning models from training data

7.1 Learning from i.i.d data

Algorithm for Bayes Net MLE:

Given BN of structure G and dataset D of complete observations

For each X_i estimate: $\hat{\theta}_{X_i|Pa_i} = \frac{\text{Count}(X_i, Pa_i)}{\text{Count}(Pa_i)}$

Pseudo-counts for lime and cherry flavor:

$$\theta_{F=c} = \frac{\text{Count}(F=c) + \alpha_c}{N + \alpha_c + \alpha_j}$$

7.1.1 Score based structure learning

Define scoring function $S(G; D)$ and search over BN structure G : $G^* = \text{argmax}_G S(G; D)$

Examples of scores:

MLE Score:

$\log P(D|\theta_G, G) = N \sum_{i=1}^n \hat{I}(X_i; Pa_i) + \text{const.}$

Where mutual information ($I(X_i, X_j) \geq 0$) is:

$$I(X_i, X_j) = \sum_{x_i, x_j} P(x_i, x_j) \log \frac{P(x_i, x_j)}{P(x_i)P(x_j)}$$

Empirical mutual information:

$$\hat{P}(x_i, x_j) = \frac{\text{Count}(x_i, x_j)}{N}$$

$$\hat{I}(X_i, X_j) = \sum_{x_i, x_j} \hat{P}(x_i, x_j) \log \frac{\hat{P}(x_i, x_j)}{\hat{P}(x_i)\hat{P}(x_j)}$$

Regularizing a Bayes Net:

$$S_{BIC}(G) = \sum_{i=1}^n \hat{I}(X_i; Pa_i) - \frac{\log N}{2N} |G|$$

where G is the number of parameters, n the number of variables and N the number of training examples.

Chow-Liu algorithm:

- For each pair X_i, X_j of variables, compute:

$$\hat{P}(x_i, x_j) = \frac{\text{Count}(x_i, x_j)}{N}$$

- Compute mutual information

- Define complete graph with weight of edge (X_i, X_j) given by the mutual information

- Find max spanning tree \rightarrow undirected tree

- Pick any variable as root and orient the edges away using breadth-first search.

8 Reinforcement Learning

8.1 Model-based RL

8.1.1 ϵ greedy

With probability ϵ , pick random action. With prob $(1 - \epsilon)$, pick best action. If sequence ϵ satisfies Robbins Monro criteria \rightarrow convergence to optimal policy with prob 1.

8.1.2 R_{max} algorithm

Input: starting x_0 , discount factor γ .

Initially: add fairy tale state x^* to MDP

- Set $r(x, a) = R_{max}$ for all states x and actions a

- Set $P(x'|x, a) = 1$ for all states x and actions a

- Choose the optimal policy for r and P

Repeat: 1. Execute policy π and, for each visited state/action pair, update $r(x, a)$

2. Estimate transition probabilities $P(x'|x, a)$

3. If observed 'enough' transitions/rewards, recompute policy π , according to current model \hat{P} and r .

'Enough'? See Hoeffding's inequality. To reduce error ϵ , need more samples N .

Theorem: With probability $1 - \delta$, R_{max} will reach an ϵ -optimal policy in a number of steps that is polynomial in $|X|, |A|, T, 1/\epsilon$ and $\log(1/\delta)$. Memory $O(|X|^2|A|)$.

8.2 Model-free RL: estimate $V^*(x)$ directly

8.2.1 Q-learning

$Q(x, a) \leftarrow (1 - \alpha_t)Q(x, a) + \alpha_t(r + \gamma \max_{a'} Q(x', a'))$

Theorem: If learning rate α_t satisfies:

$\sum_t \alpha_t = \infty$ and $\sum_t \alpha_t^2 < \infty$ (Robbins-Monro), and actions are chosen at random, then Q learning converges to optimal Q^* with probability 1.

Optimistic Q learning:

Initialize: $Q(x, a) = \frac{R_{max}}{1 - \gamma} \prod_{t=1}^{T_{init}} (1 - \alpha_t)^{-1}$

Same convergence time as with R_{max} . Memory $O(|X||A|)$. Comp: $O(|A|)$.

Parametric Q-function approximation:

$Q(x, a; \theta) = \phi(x, a)$ to scale to large state spaces. (You can use Deep NN here!)

SGD for ANNs: initialize weights. For $t = 1, 2, \dots$, pick a data point (x, y) uniformly at random. Take step in negative gradient direction. (In practise, mini-batches).

Deep Q Networks: use CNN to approx Q function. $L(\theta) = \sum_{(x, a, r, x') \in D} (r + \gamma \max_{a'} Q(x', a'; \theta) - Q(x, a; \theta))^2$

Double DQN: current network for evaluating argmax (too optimistic, and you remove θ^{old} and put θ).

8.3 Gaussian processes

A GP is an (infinite) set of random variables (RV), indexed by some set X , i.e., for each x in X , there is a RV Y_x where there exists functions $\mu : X \rightarrow \mathbb{R}$ and $K : X \times X \rightarrow \mathbb{R}$ such that for all: $A \in X$, $A = x_1, \dots, x_k$, it holds that $Y_A = [Y_{x_1}, \dots, Y_{x_k}] \sim N(\mu_A, \Sigma_{AA})$, where: Σ_{AA} = matrix with all combinations of $K(x_i, x_j)$.

K is called kernel (covariance) function (must be symmetric and pd) and μ is called mean function. **Making prediction with GPs:** Suppose $P(f) = GP(f; \mu, K)$ and we observe $y_i = f(\bar{x}_i) + \epsilon_i$, $A = \{\bar{x}_1 : \bar{x}_k\}$

$P(f(x)|\bar{x}_1 : \bar{x}_k, y_{1:k}) = GP(f; \mu', K')$. In particular, $P(f(x)|\bar{x}_1 : \bar{x}_k, y_{1:k}) = N(f(x); \mu_{x|A}, \sigma_{x|A}^2)$, where $\mu_{x|A} = \mu(\bar{x}) + \Sigma_{x,A}(\Sigma_{AA} + \sigma^2 I)^{-1} \Sigma_{x,A}^T (\bar{y}_A - \mu_A)$ and $\sigma_{x|A}^2 = K(\bar{x}, \bar{x}) - \Sigma_{x,A}(\Sigma_{AA} + \sigma^2 I)^{-1} \Sigma_{x,A}^T$.

Closed form formulas for prediction!